

## Ex01 Lih Cohen Screenshots

```
hadoop@ip-172-31-30-47:~/course
2025-11-23 18:32:40,977 INFO mapreduce.Job: map 100% reduce 0%
2025-11-23 18:32:43,991 INFO mapreduce.Job: map 100% reduce 67%
2025-11-23 18:32:44,994 INFO mapreduce.Job: map 100% reduce 100%
2025-11-23 18:32:46,003 INFO mapreduce.Job: Job job_1763921376629_0002 completed successfully
2025-11-23 18:32:46,066 INFO mapreduce.Job: Counters: 56
    File System Counters
        FILE: Number of bytes read=1188898
        FILE: Number of bytes written=6971293
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=7215145
        HDFS: Number of bytes written=889817
        HDFS: Number of read operations=42
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=6
        HDFS: Number of bytes read erasure-coded=0
    Job Counters
        Killed reduce tasks=1
        Launched map tasks=9
        Launched reduce tasks=3
        Data-local map tasks=4
        Rack-local map tasks=5
        Total time spent by all maps in occupied slots (ms)=169501920
        Total time spent by all reduces in occupied slots (ms)=75718080
        Total time spent by all map tasks (ms)=28945
        Total time spent by all reduce tasks (ms)=6465
        Total vcore-milliseconds taken by all map tasks=28945
        Total vcore-milliseconds taken by all reduce tasks=6465
        Total megabyte-milliseconds taken by all map tasks=169501920
        Total megabyte-milliseconds taken by all reduce tasks=75718080
    Map-Reduce Framework
        Map input records=138644
        Map output records=1168215
        Map output bytes=9000130
        Map output materialized bytes=1839521
        Input split bytes=1125
        Combine input records=0
        Combine output records=0
        Reduce input groups=77928
        Reduce shuffle bytes=1839521
        Reduce input records=1168215
        Reduce output records=77928
        Spilled Records=2336430
        Shuffled Maps =27
        Failed Shuffles=0
        Merged Map outputs=27
        GC time elapsed (ms)=130
        CPU time spent (ms)=12440
```

```
[hadoop@ip-172-31-30-47:~/course]
2025-11-23 18:37:04,605 INFO mapreduce.Job: map 100% reduce 100%
2025-11-23 18:37:04,610 INFO mapreduce.Job: Job job_1763921376629_0003 completed successfully
2025-11-23 18:37:04,669 INFO mapreduce.Job: Counters: 55
    File System Counters
        FILE: Number of bytes read=1236921
        FILE: Number of bytes written=8096298
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=7215145
        HDFS: Number of bytes written=889817
        HDFS: Number of read operations=57
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=12
        HDFS: Number of bytes read erasure-coded=0
    Job Counters
        Launched map tasks=9
        Launched reduce tasks=6
        Data-local map tasks=5
        Rack-local map tasks=4
        Total time spent by all maps in occupied slots (ms)=177225984
        Total time spent by all reduces in occupied slots (ms)=139384512
        Total time spent by all map tasks (ms)=30264
        Total time spent by all reduce tasks (ms)=11901
        Total vcore-milliseconds taken by all map tasks=30264
        Total vcore-milliseconds taken by all reduce tasks=11901
        Total megabyte-milliseconds taken by all map tasks=177225984
        Total megabyte-milliseconds taken by all reduce tasks=139384512
    Map-Reduce Framework
        Map input records=138644
        Map output records=1168215
        Map output bytes=9000130
        Map output materialized bytes=1930346
        Input split bytes=1125
        Combine input records=0
        Combine output records=0
        Reduce input groups=77928
        Reduce shuffle bytes=1930346
        Reduce input records=1168215
        Reduce output records=77928
        Spilled Records=2336430
        Shuffled Maps =54
        Failed Shuffles=0
        Merged Map outputs=54
        GC time elapsed (ms)=138
        CPU time spent (ms)=15300
        Physical memory (bytes) snapshot=6759374848
        Virtual memory (bytes) snapshot=135780823040
        Total committed heap usage (bytes)=7239368704
```

```
2025-11-23 19:25:50,423 INFO streaming.StreamJob: Output directory: /user/hduser/gutenberg-output
[hadoop@ip-172-31-30-47 course]$ hadoop fs -cat /user/hduser/gutenberg-output/*
the    74369
and    42768
of     34588
[hadoop@ip-172-31-30-47 course]$
```